

THEAURAX.IN · OPERATIONS DOCUMENTATION

Aurax AI Assistant

Complete Operating & Instruction Manual

WhatsApp Sales Bot · Admin Console · Knowledge Hub

Version 1.0 · 5 August 2026

Who this manual is for

The store owner and anyone on the Aurax or testing team who operates the assistant day to day. No coding knowledge is required for Parts 1–7. Part 8 onward contains technical reference material for whoever maintains the server.

Contents

1	What This System Is	
	The three parts · what the assistant can and cannot do · how a message is answered	
2	Signing In to the Admin Console	
	Addresses · team passwords · sessions · the sidebar and its badges	
3	Monitor — Watching the Bot Work	
	Every tile, chart, table and log line explained · what "normal" looks like	
4	WhatsApp — Connecting the Number	
	Pairing by QR · the linked-account panel · logging out · safe mode	
5	Knowledge Hub — Teaching the Bot	
	5.1 Teach the bot · 5.2 Knowledge sources · 5.3 Review mistakes · 5.4 The "needs answer" queue	
6	Choosing the Right Knowledge Option	
	Website crawl vs. document upload vs. Q&A · decision tree · 12 worked scenarios · mistakes to avoid	
7	Support Tickets — After-Sales	
	How tickets are raised · every column · filters · resolving and reopening	
8	What the Bot Does on Its Own	
	Instant answers · the order flow · language · follow-ups · alerts · unread backlogs and traffic bursts	
9	Your Operating Routine	
	Daily, weekly and monthly checklists · what to do after a product launch	
10	Troubleshooting	
	Symptom → cause → fix, for the fifteen problems that actually happen	
11	Technical Reference	
	Limits · settings · API endpoints · files · glossary	

1 • What This System Is

Aura is an AI sales assistant that replies to customers on WhatsApp for Theaurax.in. It answers questions, searches the live product catalogue, takes orders, sends payment links, and raises support tickets — around the clock, in English and in Tanglish.

1.1 The three parts

PART	WHAT IT DOES	WHO TOUCHES IT
The bot	Runs on the server, stays logged in to your WhatsApp number, and replies to every customer message automatically.	Nobody — it runs unattended.
The admin console	The web page you sign in to. Four sections: Monitor, WhatsApp, Knowledge Hub, Support Tickets.	You and your team, daily.
Your WooCommerce store	The real source of products, prices, stock and orders. The bot reads from it and creates real orders in it.	Managed as usual in WordPress.

1.2 What the assistant can do

- Greet customers and answer common questions instantly (COD, delivery time, sizing, returns, customization, bulk orders, tracking, jersey care, kids sizes).
- Search the live product catalogue and show matching jerseys with price and stock.
- Take a size and quantity, collect a shipping address, and confirm an order.
- Create a **real WooCommerce order** and send the customer a direct payment link.
- Look up an existing order's status for the customer who placed it.
- Raise a support ticket for wrong items, damage, delays, exchanges, or a request to speak to a human — and alert you on WhatsApp the moment it happens.
- Alert you for bulk enquiries above your configured quantity threshold.
- Re-engage a customer who went quiet mid-conversation.
- Reply in Tanglish to a customer who writes in Tanglish, for the whole conversation.

1.3 What it deliberately will not do

- It will not promise a refund, a discount, or a specific resolution to a complaint — it logs a ticket and tells the customer the team will follow up.
- It will not invent a tracking number that does not exist yet.
- It will not share an order's details with a WhatsApp number that did not place it.
- It will not answer from an uploaded document if the document does not actually cover the question — it falls back to answering normally.

1.4 How a single customer message gets answered

Understanding this order matters, because it explains **why your corrections take effect the way they do**. The bot tries the cheapest, most reliable route first and only falls through to the AI model when it has to.

Customer message arrives on WhatsApp

|



1. Your Knowledge Hub answer (Teach tab)
Strong keyword match → replied instantly, no AI

← highest priority

|

no strong match



2. Built-in FAQ (14 standard topics, EN + Tanglish)
Only when the chat is idle → replied instantly

|

not an FAQ



3. Shortcuts: "M size 2" · "2 M 5" · "yes"
Read directly as cart / confirmation, no AI

|

not a shortcut



4. AI assistant
+ your soft-matching Knowledge Hub answers
+ top 3 relevant chunks from documents / website
+ tools: search products, cart, address, order,
order lookup, support ticket, escalate

|



Reply sent back to the customer

The single most important consequence

A **Knowledge Hub answer beats everything else**, including the built-in FAQ. If the bot is saying something wrong, teaching it the right answer in the Teach tab overrides the old behaviour immediately — no restart, no developer.

2 · Signing In to the Admin Console

2.1 Where to open it

ADDRESS	WHEN TO USE IT
Your Vercel URL e.g. <code>https://aurax-admin.vercel.app</code>	The normal, everyday address. Fastest, and always up to date.
<code>http://<server>:3000/admin</code>	The backup copy served by the bot itself. Use it if the Vercel address is down. It can be slightly older than the main one.

Both work on a phone as well as a desktop. On a phone the sidebar is hidden behind the ☰ button in the top-left.

2.2 Signing in

- 1 Open the address. You will see the **AURAX Admin** sign-in card.
- 2 Type your **team password** and press Enter or click **Sign in**.
- 3 You land on the **Monitor** page. Your team name appears at the bottom of the sidebar.

There are two passwords, and which one you use decides which team the session is tagged as:

PASSWORD	SIGNS YOU IN AS	ACCESS
<code>AURAX_TEAM_PASSWORD</code>	Aurax Team	Full access to all four sections.
<code>TESTING_TEAM_PASSWORD</code>	Testing Team	Identical access — the tag exists so you know who is signed in, not to restrict anything.

Treat these passwords like the store's keys

Anyone who signs in can read every customer conversation, see phone numbers, unlink your WhatsApp number, and change what the bot says. Share them only with people who need them, and ask the developer to change them when someone leaves the team.

Why you sometimes get signed out

Sign-in sessions are held in the server's memory. Whenever the bot server restarts — a deploy, a reboot, a crash recovery — every signed-in session ends and everyone must sign in again. This is normal and is not a sign of a problem.

2.3 The sidebar

ITEM	WHAT IT IS
 Monitor LIVE	Live health of the bot: requests, tokens, active chats, API keys, server log.
 WhatsApp	Connect or disconnect the WhatsApp number the bot replies from.
 Knowledge Hub 3	Teach answers, add documents and websites, review mistakes. A red number means that many questions are waiting for you to write an answer.
 Support Tickets 2	After-sales issues. A red number is the count of open tickets awaiting your team.
Sign out	Ends your session on this device. Always use it on a shared or public computer.

Both red badges refresh every 30 seconds, on whichever page you are on.

3 · Monitor — Watching the Bot Work

The Monitor page is your instrument panel. Nothing here is a setting — it is all read-only. Statistics refresh every 5 seconds and the server log every 2.5 seconds.

3.1 The four tiles at the top

API Requests

Total AI requests made since the bot last started, split into **success** and **errors** underneath.

Read it as: a few errors alongside many successes is normal — the bot retries and switches providers automatically. Errors climbing while successes stay flat means every provider is failing; check the server log and the key list.

Tokens Used

Total text volume processed across all AI providers since restart. Tokens are how AI usage is billed and rate-limited — roughly one token per $\frac{3}{4}$ of a word.

Read it as: your cost and quota meter. A sudden jump without a matching jump in chats usually means long conversations, not more customers.

Active Chats

How many customer conversations are currently held in memory — people the bot is mid-conversation with.

Active Provider

Which AI service answered most recently, and which of its keys. You may see Groq, Sarvam, Fireworks, OpenAI, OpenRouter or Gemini here.

Read it as: switching between providers is **designed behaviour**, not a fault. The bot moves down a fallback chain whenever one provider is busy or out of quota. "None" simply means no AI call has been made since the last restart.

3.2 Call Traffic & Token Usage

The green-and-red line chart covers the **last 24 hours in eight 3-hour blocks**. Green is successful calls, red is failed ones. Under it, one bar per provider shows how much of its assumed allowance has been consumed.

What the provider bars really mean

Those percentages are measured against a built-in reference allowance per provider, not a live reading from the provider's own billing. Use them to compare providers against each other and to spot an unusual spike — not as an exact bill.

3.3 API Keys & Cooldowns

One row per key the bot holds, named `provider · Key #0`, `#1` and so on.

INDICATOR	MEANING
ACTIVE	The key is available for use right now.
COOLDOWN 42M 10S	That key hit its daily or per-minute limit. The bot has parked it and is using the others. It comes back automatically when the timer runs out — no action needed .
Green dot beside the name	This is the key that served the most recent call.

When to worry: only when *every* key of *every* provider shows a cooldown at the same time. Then customers are getting the "we're getting a lot of messages" reply, and their questions are being queued for automatic retry.

3.4 Active Chat Sessions

COLUMN	WHAT IT TELLS YOU
Customer	The customer's WhatsApp number.
State	IDLE browsing or chatting · COLLECTING_ADDRESS the bot is taking their address · CONFIRMING_ORDER one "yes" away from an order.
Language	English or Tanglish. Locked for the whole conversation on first contact.
Cart	What they are about to buy — quantity, product, size. The cart holds one product at a time.
Escalate	ALERT means this chat needs a human. You will also have had a WhatsApp alert.
Last active	How long since their last message.

The most valuable habit on this page

Scan for rows stuck in **COLLECTING_ADDRESS** or **CONFIRMING_ORDER** with a "last active" of an hour or more. That is a customer who nearly bought and stopped. The bot will nudge them automatically after 3 hours — but a personal message from you converts far better.

3.5 Server console

The dark panel at the bottom is the live server log — the last 250 lines, auto-scrolling. Grey is informational, amber is a warning, red is an error. Scroll up and it stops following so you can read; scroll back to the bottom and it resumes.

Lines worth recognising:

LINE STARTS WITH	MEANING
[Tokens] <provider> ...	A successful AI call, with its token counts. Routine.
[Retrieval Service] ...	Something about your uploaded documents or crawled website.
[Knowledge Service] ...	Something about your taught Q&A answers.
[Diagnose] ...	The half-hourly scan that finds questions the bot failed to answer.
Anything red and repeating	Copy the line and send it to your developer. That is the useful thing to report.

4 · WhatsApp — Connecting the Number

This page links a real WhatsApp account to the bot. Whatever number is linked here is the number your customers message, and every message it receives is answered by Aura.

4.1 The status badge

BADGE	MEANING	WHAT TO DO
CONNECTED & ACTIVE	The bot is live and answering customers.	Nothing.
SCAN QR CODE	Waiting for a phone to pair with it.	Scan it — see 4.2.
CONNECTING...	Starting up or reconnecting.	Wait about 30 seconds.
DISCONNECTED	Not linked to any number. Customers are not being answered.	Wait for the QR, then pair.

4.2 Pairing a number for the first time

- 1 Wait for the page to show **Scan QR code** and display the QR image.
- 2 On the phone that owns the business number, open **WhatsApp**.
- 3 Tap **Menu** (Android) or **Settings** (iPhone) → **Linked devices** → **Link a device**.
- 4 Point the phone at the QR code on screen.
- 5 The badge turns green and a ✓ replaces the QR. The bot is live.

The QR code expires

If it has been on screen a while and scanning does nothing, wait — the page refreshes every 2.5 seconds and a fresh code appears on its own. Scan the new one.

Unread chats already on the phone WILL be answered

If you connect a number with, say, 1,000 unread conversations sitting on it, the bot finds every one where the customer is still waiting for a reply and answers them — recent ones straight away, older ones gradually over the following hours.

You will get a WhatsApp summary within a minute of connecting telling you how many were found. Nothing is discarded. See section 8.7 for exactly how it is paced.

4.3 The linked-account panel

Once connected, the page shows exactly which account is paired. This exists so you never have to check "Linked devices" on every handset in the shop to find out which phone is running the bot.

FIELD	WHAT IT SHOWS
Linked number	The WhatsApp number customers are messaging, e.g. <code>+91 98765 43210</code> .
Account name	The WhatsApp profile name on that account.
Device	The phone's platform (Android, iPhone).
Linked since	How long the current session has been connected.

This panel appears **only while connected**. It is deliberately hidden when the badge is red, so a stale number can never sit next to a "Disconnected" badge and mislead you.

4.4 🗑️ Log out this number

Unlinks the phone from the bot. Use it when moving the bot to a different number, or if you suspect the session has been compromised.

This stops the bot answering customers

From the moment you confirm until someone scans a new QR code, **every incoming customer message goes unanswered**. The confirmation dialog names the exact number so you cannot unlink the wrong one by accident. Do it at a quiet hour, and have the new phone in hand.

What happens when you click it:

- A confirmation appears naming the number. Cancel is safe.
- The session is revoked on the phone **and** cleared on the server — the bot cannot silently re-pair the same number afterwards.
- After about two seconds the bot restarts its connection and a **fresh QR code** appears on the same page. No server restart is needed.
- If the unlink fails, the page says so plainly and tells you what is needed — it will not claim a clean unlink that did not happen.

4.5 Safe mode

The grey banner at the bottom of the page is a reminder that **safe mode** may be on. When the server is configured with a list of tester numbers, the bot replies **only** to those numbers and silently ignores everyone else.

The classic false alarm

"The bot is connected but customers get no reply" is almost always safe mode still switched on after testing. It is a server setting, not a button on this page — ask your developer to clear the tester-number list when you go live.

5 · Knowledge Hub — Teaching the Bot

This is the section that makes the assistant yours. Everything you save here goes live immediately — no restart, no deploy, no developer. It has three tabs.

TAB	USE IT TO...
 Teach the bot	Write exact answers to exact questions. The most direct and most reliable control you have.
 Knowledge sources	Upload documents or crawl your website so the bot can read from them.
 Review mistakes	Read real conversations that went wrong and fix them on the spot.

5.1 The "Teach the bot" tab

The form at the top adds or edits one answer. Below it is every answer you have saved.

Customer question

One realistic example of how a customer would ask. This is a **label for you** — it helps you find the entry later, and it is what appears in the saved list. It is not what the matching is based on.

Example: *Do you deliver to Sri Lanka?*

Correct answer

Word for word what the bot should reply. When the match is strong, this text is sent to the customer exactly as written. Write it in the tone you want customers to hear — complete sentences, your own phrasing, emoji if you like.

Keep it to a WhatsApp-sized reply. Two or three sentences beats a paragraph.

Trigger keywords — comma separated

This is the field that actually decides when your answer is used. The question field does not; the keywords do. Getting these right is the whole skill.

Two kinds of keyword behave differently:

- **Phrases** (two or more words, e.g. `sri lanka`) are a **strong** signal. One phrase match is enough for the bot to reply with your answer directly.
- **Single words** (e.g. `international`) are a **weaker** signal. It takes **two** different single words matching for a direct reply. A single word on its own only nudges the AI toward your answer.

Practical rule: give every entry at least one phrase, plus three or four single words.

Language

SETTING	USE WHEN
Both (default)	Almost always. The answer is available to every customer.
English only	The answer is written in English and would read oddly to a Tanglish customer.
Tanglish only	The answer is written in Tanglish.

To answer the same question properly in both languages, create **two entries** — one English-only, one Tanglish-only, each with keywords in its own language. The bot picks the one matching the customer's locked language.

Active checkbox

Ticked, the bot uses this answer. Unticked, it is saved but ignored — useful for a seasonal answer (a festival offer, a stock pause) you want to switch back on later rather than retype.

 **Save answer** stores it and it is live on the very next customer message. Saving with an empty answer is refused.

Reading the saved-answers list

CHIP ON THE CARD	MEANING
NEEDS ANSWER	The system found a question the bot failed on and drafted it for you. It is not being used until you write an answer.
ASKED 4×	How many times that unanswered question has come up. Answer the high numbers first.
FROM A REAL CHAT	You created this from the Review tab.
AUTO-FOUND	Started as an auto-draft and you have since answered it.
BOTH / ENGLISH / TANGLISH	The entry's language setting.
INACTIVE	Saved but switched off.

Edit reopens the entry in the form at the top. **Delete** removes it permanently (with a confirmation) and the bot stops using it at once.

5.2 The "📖 Knowledge sources" tab

Where you give the bot material to read: your website, or files you upload. Part 6 covers *which* to use; this section covers *how*.

Website

FIELD	WHAT TO ENTER
URL	The page to start from. Point it at the most specific page that holds the information — <code>https://theaurax.in/shipping-policy</code> , not just <code>https://theaurax.in</code> .
Max pages default 15	How many pages the crawl may read before stopping. Keeps a crawl from running away.
Link depth default 2	How many clicks away from the starting page it may travel. 0 = that one page only. 1 = it and the pages it links to. 2 = one step further.

 **Crawl & index** starts it. The button reads "Crawling..." while it works — this can take a minute or two. When it finishes you get a message like "*Indexed 6 page(s) — 41 chunks.*"

Depth 2 from the homepage is usually a waste

A real crawl of theaurax.in at 15 pages / depth 2 spent its entire budget on product-category grids and empty account pages, and **never reached the shipping, returns or about pages** — the exact prose the feature exists to index. The crawler follows links in alphabetical order, so category pages consume the budget first.

Do this instead: crawl your policy pages one at a time at **depth 0**. Three precise crawls beat one broad one every time.

Re-crawling the same URL does not replace the old copy

It creates a **second** source, and the bot then searches both — near-identical duplicates competing with each other. **Always delete the existing source first**, then crawl the URL again.

Document

Choose a file and click  **Upload & index.**

REQUIREMENT	DETAIL
Accepted types	PDF, DOCX, TXT, MD, HTML
Maximum size	20 MB per file
Must contain real text	A scanned or photographed PDF will be rejected . If you can select and copy the text in a PDF reader, it will work; if it is a picture of a page, it will not. Re-export it from Word, or retype it into a TXT file.

Uploaded files are read into memory, indexed, and never written to the server's disk.

Reading a source card

CHIP	MEANING
WEBSITE / DOCUMENT	Which kind of source it is.
41 CHUNKS	How many searchable pieces it was split into. Roughly one per 900 characters.
118.1 KB	How much text was extracted.
6 PAGES	Pages crawled, or pages in the PDF.
SEMANTIC SEARCH	Good. The bot understands re-worded questions against this source.
KEYWORD ONLY	The bot can only match this source on the exact words used. It will miss paraphrases.

Turn off / Turn on takes a source out of use immediately without deleting it — the fastest way to test whether a source is causing bad answers. **Delete** removes the source and all of its chunks permanently.

The **Indexed knowledge** bar shows total size across all sources. There is no hard limit; the bar is there to give you a sense of scale.

If a "⚠️ Keyword-only matching" warning appears at the top of the tab

Semantic search is temporarily unavailable, so anything indexed *while it is showing* will only ever match on exact words. Wait until it clears, then delete and re-add those sources so they pick up proper semantic indexing.

5.3 The "🔍 Review mistakes" tab

A list of real conversations that show a sign of trouble. A chat is flagged when it contains:

- an **error or fallback reply** — the bot said it couldn't process something;
- a **repeated question** — the customer asked the same thing twice or more, which almost always means the first answer missed;
- an **abandoned cart** — they got partway through buying and stopped.

Each card shows the customer, the reasons it was flagged, and the last few turns of the real conversation. Read them — this is the single best source of ideas for what to teach.

✏️ **Teach the right answer** jumps you to the Teach tab with the customer's actual question already filled in and keywords suggested from it. Write the answer, check the keywords, save. **This is the core improvement loop of the whole system.**

5.4 The "needs answer" queue and the red badge

Every 30 minutes the server scans conversations for questions the bot stumbled on, and drafts each one as an **inactive** entry at the top of your Teach tab. The red number on the sidebar is how many are waiting. You also get a WhatsApp ping when new ones appear.

Two ways to clear one:

ACTION	USE WHEN	RESULT
 Answer	It is a real question worth answering.	Write the reply and save. It becomes a live answer and the badge drops by one.
Dismiss	It is noise — a typo, a wrong number, something irrelevant.	Permanently removed. It will never be suggested again , even though the old conversation still exists.

Use Dismiss freely

Dismiss exists precisely so the queue stays meaningful. A queue full of junk you keep scrolling past is a queue you will stop reading. Ten minutes a week keeping it at zero is the highest-value maintenance you can do on this system.

6 · Choosing the Right Knowledge Option

Three ways to give the bot information, and picking the wrong one is the most common reason people conclude "it doesn't work". This part exists to make the choice obvious.

6.1 The three options at a glance

Teach (Q&A)

Exact answers, word for word

- Highest priority — beats everything
- Replies instantly, no AI involved
- You control the exact wording
- One question at a time
- Live the moment you save

Document

Reference material to read from

- Feeds the AI, doesn't override it
- Handles many facts in one upload
- Wording is the AI's, not yours
- Needs selectable text
- Best for tables and lists


Website

Your live pages, indexed

- Same behaviour as a document
- No file to prepare
- Only as good as the page's text
- Does **not** auto-refresh
- Best for existing policy pages

6.2 The decision tree

START: what do you want the bot to know?



6.3 When to use each — the full table

YOU WANT THE BOT TO...	USE	WHY
Quote the exact COD charge	Teach	A number must be exact. Only Teach guarantees the words sent.
Give the correct returns window	Teach	A policy promise. Never leave the wording to the AI.
Answer "do you ship to Sri Lanka?"	Teach	A yes/no with consequences. One sentence, said the same way every time.
Share the wholesale contact number	Teach	Digits must be perfect.
Announce a festival offer for two weeks	Teach (untick Active after)	Trivial to switch off when it ends.
Fix a wrong reply seen in a real chat	Review → Teach	Pre-fills the customer's real words for you.
Know a 40-row size chart	Document	Far too many facts for one Q&A. Chunking keeps a table intact.
Know a 6-page club-wise price list	Document	Many facts, structured, changes rarely.
Explain a long customization process	Document	The AI can summarise the relevant part for each customer.
Know the fabric / GSM / care details	Document	Detail-heavy reference nobody wants to retype as Q&As.
Use your existing shipping-policy page	Website (depth 0)	Already written and public — no file to make.
Use your About / brand story page	Website (depth 0)	Background prose, no exact wording needed.
Use your FAQ page	Website (depth 0)	One crawl covers many questions at once.
Know current prices and stock	None — it already does	Live from WooCommerce. Never put prices in a document; it will go stale and conflict.

6.4 Website crawl or document upload? The five-question test

QUESTION	YES →	NO →
Is the information already on a public page of your site?	Website	Document
Is that page mostly readable paragraphs, not a grid of product tiles?	Website	Document
Will the page's wording change often?	Document*	Website
Is it a table, chart, or list of measurements?	Document	Either
Is it internal — not published on the website at all?	Document	Website

* Neither option updates itself. But a document you control is easier to re-upload deliberately than a page that quietly changed under you without anyone re-crawling it.

Neither documents nor website crawls refresh automatically

An indexed source is a **snapshot taken the moment you added it**. Edit the page on your website, and the bot keeps answering from the old text until someone deletes the source and crawls it again. Put a reminder in your calendar whenever you change a policy page.

6.5 Twelve real scenarios

SITUATION	DO THIS	NOTES
"COD charges keep being answered wrong."	Teach	Keywords: <code>cod charge, cash on delivery, cod fee, cod, charges</code>
"We just started shipping to Nepal."	Teach	Add the entry; also delete any older entry that said India-only.
"Customers keep asking for the size chart."	Document	Upload the chart as a real-text PDF or DOCX.
"Our shipping policy page is already detailed."	Website , depth 0	Crawl that exact URL, not the homepage.
"We changed the returns window from 7 to 10 days."	Teach + re-crawl	Teach it first so it is right today, then delete and re-crawl the policy page.
"The bot mixes up two similar club jerseys."	Neither	That is product search, not knowledge. Fix the product names/tags in WooCommerce and re-sync.
"It quotes an old price."	Delete the source	A price got indexed from a document or old page. Remove it — live prices come from WooCommerce.
"We want a Tanglish reply for delivery time."	Teach ×2	One entry Tanglish-only with Tanglish keywords, one English-only.
"Diwali offer for two weeks."	Teach	Untick Active when it ends instead of deleting.
"Customers ask about washing the jersey."	Check first	Jersey care is already a built-in FAQ. Only teach an entry if the built-in answer is wrong for you.
"We have a 30-page brand deck."	Neither	Marketing decks are slide fragments — they index badly. Extract the 5 useful facts into Q&As.
"Bot said something we can't stand behind."	Review → Teach , immediately	Highest priority. Your answer overrides everything from the next message onward.

6.6 Mistakes to avoid

MISTAKE	WHAT GOES WRONG
Crawling the homepage at depth 2 and expecting policies	The budget is eaten by category grids and account pages. Crawl the policy page itself at depth 0.
Re-crawling a URL without deleting the old source	Two near-identical copies compete in search. Delete first, always.
Uploading a scanned or photographed PDF	Rejected — there is no text in it, only an image. Re-export or retype.
Putting prices or stock in a document	It will contradict the live catalogue the day something changes.
One keyword per Q&A entry	A single word only nudges the AI. Give a phrase plus three or four words.
Very generic keywords like <code>jersey</code> or <code>order</code>	Fires on almost every message and hijacks unrelated conversations.
Trusting a document for a promise you must honour	The AI paraphrases documents. Promises belong in the Teach tab.
Uploading everything "just in case"	More weak material makes it harder for the right chunk to win. Add deliberately.
Never opening the Review tab	The system finds your problems for you, but only if someone reads them.

If you remember one paragraph from this manual

Anything you must say exactly — a price, a policy, a promise, a phone number — goes in the Teach tab.

Documents and website crawls are for background knowledge the AI may summarise in its own words. When in doubt, teach it.

7 · Support Tickets — After-Sales

When a customer has a problem the bot cannot solve, it raises a ticket and alerts you on WhatsApp straight away. This page is the shared list your team works from. It refreshes every 10 seconds.

7.1 How a ticket gets created

The bot raises one on its own when a customer reports an after-sales issue. Before doing so it apologises, collects their **name**, the **order ID** and a **short description**, and asks for a **photo** when the issue is a wrong, damaged or misprinted item. Then it:

- creates the ticket with a reference like `TKT-1043` ;
- sends you the full details on WhatsApp immediately, photo included;
- gives the customer the reference number and the support email;
- promises only that the team will follow up — **never** a refund or a specific outcome.

7.2 The four tiles

TILE	MEANING
 Total Tickets	Every ticket ever raised.
 Open	Awaiting your team. This is the number on the sidebar badge.
 Resolved	Closed out.
 With Photos	Tickets where the customer sent an image — already forwarded to you on WhatsApp.

7.3 The ticket table

COLUMN	WHAT IT SHOWS
Ticket	The reference the customer was given. Use it when you call them back.
Customer	Name, WhatsApp number, and email if collected.
Issue	The category — see below.
Order	The order number, or — if the customer could not give one.
Photo	 means an image came with it, already sent to your WhatsApp.
Details	The problem in the customer's own words.
Created	How long ago it was raised.
Action	Resolve , or Reopen if already resolved.

Issue types

LABEL	MEANING	PRIORITY
WRONG ITEM	They received a different product.	Urgent
DAMAGED	Arrived torn, stained or defective.	Urgent
MISSING PACKAGE	Marked delivered but not received.	Urgent
DELAYED DELIVERY	Taking longer than expected.	Same day
WRONG CUSTOMIZATION	Wrong name or number printed.	Same day
EXCHANGE	Size exchange request.	Same day
TALK TO HUMAN	Explicitly asked for a person.	Fast — they are already impatient
OTHER	Anything else.	Read it

7.4 Working the list

- **Open / Resolved / All** switches the view. It opens on **Open** — your work queue.
- The **search box** filters by name, phone, ticket ID, order number, issue type or any word in the description. Searching the order number is the fast way to find a returning customer's history.

- **Resolve** marks it done and removes it from the Open list and the sidebar badge. **Mark it resolved only after you have actually contacted the customer** — the bot does not tell them anything when you click it.
- **Reopen** puts a resolved ticket back in the queue if the customer comes back about it.

A working rhythm

Clear **Wrong item**, **Damaged** and **Missing package** first — they cost you a customer if they sit. Then **Talk to human**: that customer is already frustrated and the bot has told them a person is coming. Aim to end each day with the sidebar badge at zero.

8 · What the Bot Does on Its Own

Nothing in this part needs managing. It is here so that when a customer forwards you a screenshot, you know exactly why the bot said what it said.

8.1 Instant answers

Four kinds of message are handled in code, with no AI call at all — which is why they are fast, consistent and never invent anything:

MESSAGE	HANDLED AS
A question matching one of your taught answers	Your answer, sent word for word.
A standard question while the chat is idle	The built-in FAQ — 14 topics, English and Tanglish.
"M size 2", "1st one, L 3", "2 M 5", "XL"	Read directly as product, size and quantity, and put in the cart.
"yes", "confirm", "seri", "ok" while confirming	The order is placed immediately.

The confirmation shortcut only fires when the message is *entirely* a confirmation. "Yes but change the address" correctly goes to the AI instead.

8.2 The order flow

```
Customer asks for a jersey
  ↓ bot searches the live WooCommerce catalogue
Bot shows matching products with price and stock
  ↓ customer gives size and quantity
Cart updated                                [state: IDLE]
  ↓ bot asks for the shipping address
Address collected                          [state: COLLECTING_ADDRESS]
  ↓ bot reads the full order back
Customer confirms                          [state: CONFIRMING_ORDER]
  ↓
A REAL WooCommerce order is created
  ↓
Customer receives a direct payment link    [state: IDLE]
```

The order appears in WooCommerce exactly like any other. If order creation fails, the bot falls back to sending a branded PDF proforma invoice so the sale is not lost. The cart holds one product at a time — a new product

replaces the previous one.

8.3 Language

The language is decided from the customer's **first** message and locked for the whole conversation, so the bot never drifts between English and Tanglish mid-chat. Tanglish conversations are routed to an AI provider trained on code-mixed Indian text; English conversations go to a faster general provider.

8.4 Cold-lead follow-up

Every 30 minutes the bot looks for customers who went quiet. Anyone inactive for **3 hours or more** gets a friendly re-engagement message, mentioning what was in their cart if anything was. **Maximum two follow-ups per customer** — it will not pester.

8.5 When you get pinged on WhatsApp

ALERT	WHY
 New Support Ticket	An after-sales issue was logged. Full details and photo included.
Bulk order enquiry	Someone asked about a quantity at or above your bulk threshold. Worth a personal reply.
Escalation	The bot judged that a human should take this conversation over.
Knowledge gap found	New questions the bot could not answer are waiting in the Teach tab.

8.6 When every AI provider is busy

The bot works down a chain of providers and keys automatically. If they are all exhausted at once, the customer gets a polite "we're getting a lot of messages" reply and **their question is saved to a retry queue**. It is answered automatically once quota frees up — and the queue survives a server restart, so nothing is silently dropped.

8.7 Missed messages — nobody is left unanswered

The bot only sees messages that arrive while it is actually connected. On its own, that would mean two groups of customers never got a reply: people who were already sitting unread on the phone when you paired it, and anyone who wrote while the server was restarting or down.

A catch-up sweep closes both gaps. Every time the bot connects, it checks every chat and finds the ones where the customer spoke last and never got an answer — that is exactly the definition of someone you owe a reply to. Those customers are then answered.

SITUATION	DOES THE BOT REPLY?
A message arrives while the bot is connected and running	YES — within seconds.
Unread chats already on the phone when you paired the bot	YES — found by the sweep and answered.
Messages that arrived while the server was down or disconnected	YES — found the moment it comes back.
A chat where you or a staff member already replied from the phone	NO — correctly left alone. The bot only answers where nobody has.
A message received but not answerable (all AI providers busy)	YES — queued and retried automatically. See 8.6.

How the catch-up is paced

Answering a thousand-message backlog as fast as possible is the single most reliable way to get a WhatsApp number banned. So missed customers are split into two groups:

GROUP	WHEN THEY WROTE	WHAT HAPPENS
Still waiting	Within the last 12 hours	Answered right away when the bot reconnects. These people are very likely still interested.
Older backlog	More than 12 hours ago	Answered gradually — roughly two per minute, and only while no live customer is mid-conversation. A 1,000-chat backlog clears in about eight hours.

Three things that make this safe rather than embarrassing

- **Old messages are answered as old messages.** Anything more than a day late opens by apologising for the delay, rather than replying as if it just arrived.
- **Live customers always come first.** The backlog pauses itself whenever someone is actually chatting, and resumes when the line is clear.
- **Nobody is answered twice.** If a queued customer messages again before their turn comes, the old queued reply is dropped and they are handled as a normal live conversation.

It survives restarts

The queue of missed customers is written to the database, not just held in memory. If the server restarts halfway through clearing a backlog, it picks up exactly where it left off — which is the whole point of the feature.

You will be told. Within a minute of the bot connecting you get a WhatsApp message like: *"Found 47 customers waiting for a reply — 6 recent, being answered now; 41 older, going out gradually over about 1 hour."* A second message confirms when the backlog is fully cleared.

The one thing still worth doing yourself

A very old backlog is answered, but slowly and by a bot. If there are **high-value or angry customers** in there, message them personally rather than waiting for their turn in the queue. The catch-up guarantees nobody is ignored — it does not replace judgement about who deserves a call.

8.8 What happens when many customers message at once

Every customer is processed in parallel, so thirty people messaging together are all understood at the same time — nobody is stuck behind anybody else, and nothing is dropped. The limit is on **sending**, not thinking.

All outgoing messages share one paced queue, because WhatsApp measures behaviour per account, not per chat. Replying to everyone instantly is the clearest possible automation signal and is what gets numbers banned. So replies leave at roughly one every 1.7 seconds.

CUSTOMERS MESSAGING AT ONCE	EVERYONE ANSWERED WITHIN	THE LAST PERSON WAITS
30	about 51 seconds	about 51 seconds
90	about 2 minutes 49 seconds	about 2 minutes 49 seconds

Measured against the live send pacer with the current settings. Roughly 32–35 replies per minute either way.

Customers see "typing..." the whole time they are waiting

The typing indicator starts when their message arrives and only stops when the reply actually goes out. A queued customer sees an active, busy shop rather than silence — which is why a 50-second wait during a rush is far less damaging than it sounds.

Where it starts to show: around 30 simultaneous customers the delay is barely noticeable. A reel drop putting 200 or more people on you at once is where the queue becomes visible. Photos cost double, because forwarding the image to you uses a slot as well as the reply, and ticket and escalation alerts share the same queue.

If large bursts become routine, the pacing can be raised — but that trades directly against ban risk, so it is a decision for your developer with real evidence from a real spike, not a default to change.

9 · Your Operating Routine

9.1 Every day — about 5 minutes

- 1 **WhatsApp page** — confirm the badge is green and the linked number is the right one.
- 2 **Tickets** — clear the open queue. Urgent categories first.
- 3 **Knowledge Hub badge** — if there is a red number, answer or dismiss those questions.
- 4 **Monitor** — glance at the traffic chart. Green rising, red flat is healthy.

9.2 Every week — about 20 minutes

- 1 **Review mistakes tab** — read every flagged conversation. Teach an answer for each real gap.
- 2 **Saved answers** — skim your entries. Delete anything that is no longer true.
- 3 **Active chats** — look for stalled carts and message those customers personally.
- 4 **Knowledge sources** — check nothing has drifted out of date since you indexed it.

9.3 Every month — about 30 minutes

- Re-index any source whose underlying page or file has changed. **Delete first, then re-add.**
- Check every source still shows **SEMANTIC SEARCH**. If one says keyword only, delete and re-add it.
- Ask your developer to run the product sync so the catalogue matches WooCommerce.
- Read a handful of ordinary, unflagged conversations end to end. Flags catch failures; only reading catches an answer that is technically fine but sounds wrong for your brand.

9.4 Before and after a product launch or reel drop

WHEN	DO THIS
The day before	Run the product sync so new items are searchable. Teach a Q&A for the launch itself ("when does the new kit drop?"). Confirm WhatsApp is connected.
During	Keep the Monitor page open. Expect provider switching and key cooldowns under a traffic spike — that is the system working, not failing.
The day after	Review flagged conversations while the traffic is fresh — a launch produces the richest batch of real questions you will ever get. Teach them all.

10 · Troubleshooting

SYMPTOM	MOST LIKELY CAUSE	WHAT TO DO
Customers get no reply at all	WhatsApp shows Disconnected, or safe mode is still on.	Open the WhatsApp page. If red, wait for the QR and re-pair. If green but still silent, safe mode is on — ask your developer to clear the tester-number list.
Only some people get replies	Safe mode, almost certainly.	The numbers that work are the tester list. Have it cleared.
The old backlog is being answered slowly	Not a fault. Older missed customers are deliberately dripped out to protect the number from a ban.	Expected — roughly two per minute. Message any high-value customers personally instead of waiting for their turn. See section 8.7.
Replies are slow during a rush	The send queue is pacing them to protect the number from a ban.	Expected. Thirty simultaneous customers all get answered inside about a minute, and they see "typing..." while they wait. Only worth changing if 200+ bursts are routine.
No "missed messages" summary after a restart	Either nothing was missed, or catch-up is switched off on the server.	No summary means no missed customers, which is good news. If you expect one and it never comes, ask your developer to confirm catch-up is enabled.
A customer got a reply to a very old message	Working as designed — the backlog sweep reached them.	Nothing to fix. The reply opens by apologising for the delay.
"We're getting a lot of messages" replies	Every AI key is in cooldown.	Check the key list on Monitor. It recovers on its own; those questions are queued and answered automatically. If it happens daily, you need more capacity — tell your developer.
The bot ignores an answer I taught	Keywords too weak, or wrong language setting.	Add a two-word phrase keyword plus three or four single words. Confirm Active is ticked and the language is Both .
A taught answer fires on the wrong questions	Keywords too generic.	Remove broad words like <code>jersey</code> , <code>order</code> , <code>price</code> . Keep distinctive phrases only.
Uploading a PDF fails	It is a scan, over 20 MB, or an unsupported type.	Try selecting text in a PDF reader — if you can't, it's an image. Re-export from Word or retype into a TXT file.

Crawl says the site is blocking us	A security plugin, Cloudflare rule, or maintenance mode.	The on-screen message names the cause. It is fixed in WordPress, not here. Maintenance mode is a toggle in wp-admin.
Crawl finished but the bot still doesn't know it	The crawl never reached the page with the information.	Check the source card's page count. Delete it and crawl the specific policy URL at depth 0.
A source says "keyword only"	Semantic search was unavailable when it was indexed.	Once the warning at the top of the tab has cleared, delete and re-add the source.
The bot gives outdated information	An old source or a stale taught answer.	Search your saved answers first and fix or delete. Then check sources — turn one off to test whether it is the culprit.
Prices quoted are wrong	A price was indexed from a document or old page.	Delete that source. Live prices come from WooCommerce and should never be duplicated into knowledge.
Signed out unexpectedly	The server restarted.	Sign in again. Normal.
Console won't load at all	The bot server is down, or the browser cached an old version.	Try the backup address <code>http://<server>:3000/admin</code> . If that also fails, the server is down — contact your developer.
Red badge count never drops	Drafts are being answered but new ones keep arriving.	Healthy, if the questions are genuinely new. Dismiss the junk so the number stays meaningful.
Bot mixes up two similar jerseys	Product naming in WooCommerce, not knowledge.	Make the product titles and tags more distinct in WooCommerce and re-sync. No amount of teaching fixes catalogue search.

What to send your developer

The **exact customer message**, the **bot's exact reply**, the **time it happened**, and any **red lines from the server console** around that time. Those four things solve most reports in one pass; "it's not working" does not.

11 • Technical Reference

11.1 Limits and defaults

ITEM	VALUE
Document upload — maximum size	20 MB per file
Document upload — accepted types	.pdf, .docx, .txt, .md, .html, .htm
Website crawl — default / maximum pages	15 / 100
Website crawl — default / maximum depth	2 / 3
Website crawl — scope	Same origin only, 400 ms delay between pages
Text chunk size / overlap	≈900 characters / 150 characters
Chunks injected per customer question	Top 3, capped at 2,400 characters
Relevance threshold for injection	0.33 (hybrid: 75% meaning, 25% keyword)
Q&A direct-answer threshold	One phrase keyword, or two single-word keywords
Built-in FAQ topics	14, English and Tanglish
Cold-lead follow-up	After 3 hours idle, maximum 2 per customer, checked every 30 min
Unanswered-question scan	20 s after start, then every 30 minutes
Server console buffer	Last 250 lines
Refresh rates	Stats 5 s · logs 2.5 s · WhatsApp 2.5 s · tickets 10 s · badges 30 s
Outbound WhatsApp pacing	≥1.2 s between sends plus random jitter, max 30/minute (ban protection)
Measured burst throughput	32–35 replies/minute · 30 customers ≈ 51 s · 90 customers ≈ 2 min 49 s
Missed-message catch-up — recent	Within 12 hours: answered immediately on reconnect
Missed-message catch-up — older backlog	≈120/hour, no age limit (≈1,000 chats in 8 hours)
Duplicate-message protection	Last 1,000 message IDs remembered, so a resend is never answered twice

11.2 Settings that live on the server

These are not editable from the console — they are set in `apps/bot/.env` by whoever maintains the server.

SETTING	CONTROLS
<code>AURAX_TEAM_PASSWORD</code>	Aurax team console password
<code>TESTING_TEAM_PASSWORD</code>	Testing team console password
<code>OWNER_WHATSAPP_NUMBER</code>	Where ticket and escalation alerts are sent
<code>BULK_ORDER_THRESHOLD</code>	Quantity that triggers a bulk-order alert
<code>ALLOWED_TEST_NUMBERS</code>	Safe mode — if set, only these numbers get replies
<code>SUPPORT_EMAIL</code> / <code>WHOLESALE_NUMBER</code>	Contact details the bot hands out
<code>WOOCOMMERCE_*</code>	Connection to the store: products, stock, orders
<code>ADMIN_ALLOWED_ORIGINS</code>	Which console addresses may talk to the server
<code>MONGODB_URI</code>	Database. Without it, data is stored in local files
<code>CATCHUP_ENABLED</code>	Missed-message catch-up. Leave on — off means missed customers are never answered
<code>CATCHUP_FRESH_HOURS</code>	How recent counts as "still waiting" (default 12)
<code>CATCHUP_DRAIN_PER_HOUR</code>	Backlog drip rate. Raising it raises ban risk

11.3 API endpoints behind the console

Every endpoint except login requires the bearer token issued at sign-in.

ENDPOINT	PURPOSE
POST /api/knowledge-hub/login	Sign in, returns a token
GET /api/whatsapp/status	Connection state, QR, linked account
POST /api/whatsapp/logout	Unlink the paired number
GET /api/provider-stats	Monitor tiles, chart, key cooldowns
GET /api/sessions	Active chat sessions table
GET /api/logs	Server console buffer
GET /api/retry-stats	Retry queue and provider status
GET / POST /api/knowledge	List and save taught answers
DELETE /api/knowledge/:id	Delete a taught answer
POST /api/knowledge/:id/dismiss	Permanently dismiss an auto-draft
GET /api/knowledge/pending-count	Sidebar red badge
POST /api/knowledge/diagnose	Re-scan for unanswered questions
GET /api/knowledge/review	Flagged conversations
GET /api/knowledge/sources	List indexed sources
POST /api/knowledge/sources/document	Upload and index a file
POST /api/knowledge/sources/website	Crawl and index a URL
POST /api/knowledge/sources/:id/toggle	Switch a source on or off
DELETE /api/knowledge/sources/:id	Delete a source and its chunks
GET /api/tickets	All tickets
GET /api/tickets/open-count	Sidebar red badge
POST /api/tickets/:id/status	Resolve or reopen

11.4 The tools the AI can call

TOOL	WHAT IT DOES
<code>search_products</code>	Searches the live catalogue
<code>update_cart</code>	Sets product, size and quantity
<code>set_shipping_address</code>	Records the delivery address
<code>confirm_order</code>	Creates the real WooCommerce order and payment link
<code>lookup_order</code>	Reads an order's status, only for the number that placed it
<code>create_support_ticket</code>	Logs an after-sales ticket and alerts you
<code>escalate_to_human</code>	Flags the chat for a person and alerts you

11.5 Commands for whoever maintains the server

COMMAND	PURPOSE
<code>npm run bot</code>	Start the bot and API
<code>npm run sync</code>	Refresh the product catalogue from WooCommerce
<code>npm run review</code>	Flag problem conversations from the command line
<code>npm run test-agent</code>	Run the AI agent test suite
<code>npm run test-embeddings</code>	28-check retrieval regression suite
<code>npm run build-admin</code>	Rebuild the console copy the bot serves at <code>/admin</code>

11.6 Glossary

TERM	PLAIN MEANING
Chunk	A ~900-character slice of a document or page. Chunks are what the bot actually searches.
Semantic search	Matching by meaning, so a re-worded question still finds the right passage.
Keyword only	Matching by exact words. Works, but misses paraphrases.
Token	The unit AI usage is measured in — roughly $\frac{3}{4}$ of a word.
Provider	The AI service answering. The bot uses several and switches automatically.
Cooldown	A key resting after hitting its limit. Temporary and self-healing.
Fallback chain	The ordered list of providers tried until one succeeds.
Session / state	One customer's ongoing conversation and where it has reached.
Tanglish	Tamil written in English letters, mixed with English.
Safe mode	A setting that limits replies to a list of test numbers.
Escalation	The bot handing a conversation to a human and alerting you.
Draft / needs answer	A question the bot failed on, queued for you to answer. Never shown to customers.

The three habits that keep this system good

1. Clear the tickets queue daily. **2.** Read the Review tab weekly and teach an answer for every real gap. **3.** Put anything you must say exactly into the Teach tab, never into a document.

Everything else in this manual is detail. Those three keep the assistant getting better instead of quietly getting worse.