

Where to run the bot 24/7 for ~1,000 customers a day

Neenga kudutha 10 options-a, ippo run aagura **real stack**-oda constraints vachi filter panniruken — WhatsApp Web + Puppeteer (headless Chrome) + MongoDB. Ella cost-um INR/month, 1,000 conversations/day load-ku sizing pannirukken.

Stack: Node.js · whatsapp-web.js · Puppeteer · MongoDB **Load:** ~1,000 convos/day

Runtime: 24×7 always-on

THE SHORT ANSWER

Three picks, depending on what you optimise for

BEST VALUE OVERALL

Hetzner CX22

≈ ₹430 / month

2 vCPU / 4GB RAM — the exact size Chrome + Mongo wants, at the lowest honest price. EU datacenter, but that doesn't hurt you (LLM is US-based Groq anyway). Full root, PM2, persistent disk.

BEST FOR INDIA / UPI

DigitalOcean 4GB

≈ ₹2,000 / month

Bangalore datacenter, lowest India latency, best docs. Costs ~4× Hetzner but buys you local DC + the smoothest setup experience. Use its ₹100/mo backups.

CHEAPEST INDIA + INR BILLING

Indian VPS 4GB

≈ ₹600–900 / month

MilesWeb / AIC-class 4GB KVM. Direct UPI billing, Mumbai DC. Fine technically — just insist on **4GB** and NVMe; skip the ₹99 1GB teasers.

STEP 1 – WHAT THIS APP ACTUALLY DEMANDS

Your bot isn't a normal web app, so the ranking changes

A WhatsApp-Web bot behaves differently from a stateless website. Four constraints decide everything — and two of them quietly disqualify the "easy git-push" platforms.

Persistent disk for the WhatsApp session

CRITICAL

The logged-in session lives on disk in `.wwwebjs_auth/`. If the filesystem is wiped on every deploy (the default on Render / Railway / DO App Platform), you re-scan the QR code **every single deploy**. A raw VPS keeps this for free.

Always-on, single instance — no autoscaling

CRITICAL

One WhatsApp session = one browser. You **cannot** run 2 copies (double logins fight). So platforms built around horizontal autoscaling are the wrong shape, and any plan that "sleeps" on idle will drop your WhatsApp connection. You need exactly one box that never sleeps.

4GB RAM for Chrome + Mongo

HIGH

Headless Chromium (~400–600MB) + Node + a local MongoDB comfortably wants **4GB**. The ₹350 / 1GB plans will OOM-kill Chrome under real traffic. 2GB is the risky floor; 4GB is the safe target.

India-datacenter latency

LOW
PRIORITY

Tempting to over-weight this, but your AI inference is **external** (Groq is US-hosted) and WhatsApp uses a global edge. A Hetzner box in Germany answers customers just as fast. India DC only matters for local support & INR billing convenience — not speed.

STEP 2 – SIZING THE 1,000-CUSTOMERS/DAY LOAD

What 1,000 conversations really costs the server

The heavy lifting (LLM inference) is offloaded to Groq/Sarvam/OpenAI, so the server itself only juggles Chrome, the message queue, and MongoDB. The load is steady, not spiky — well within a single 4GB box.

~13k MESSAGES / DAY (IN + OUT)	~3k LLM CALLS / DAY (≈% DETERMINISTIC)
~40/min PEAK-HOUR INBOUND	4GB RAM TARGET

Math: ~1,000 convos × 6–7 messages each ≈ 13k messages/day. The bot's deterministic fast-paths (FAQ, size/qty parsing, order-confirm) skip the LLM on roughly half the turns, so only ~3k calls actually hit an API — and those don't tax *your* CPU. The built-in per-sender queue (concurrency 5) absorbs the ~40/min peak easily. **Verdict: 2 vCPU / 4GB is right; 8GB is only worth it if you later co-host Mongo Atlas-style analytics on the same box.**

STEP 3 – ALL 10 OPTIONS, SCORED FOR YOUR STACK

Fit rating = how well it matches the 4 constraints above

PROVIDER	TYPE	SUGGESTED PLAN	INR / MO*	FIT	VERDICT
Hetzner Cloud EU / US DC	Unmanaged VPS	CX22 · 2vCPU / 4GB / 40GB	₹430	A+ · 10/10	Top pick Best price-to-spec anywhere. Persistent disk, root, no sleep.
DigitalOcean Bangalore DC	Unmanaged VPS	Droplet · 2vCPU / 4GB	₹2,000	A · 9/10	India + easy Local DC, best docs. Pay ~4× Hetzner for the polish.
Vultr Mumbai /	Unmanaged VPS	HF · 2vCPU /	₹1,650	A · 9/10	India + fast CPU High-clock cores,

Delhi DC		4GB			Indian DCs. Watch bandwidth caps per region.
Linode / Akamai Mumbai DC	Unmanaged VPS	Shared · 2vCPU / 4GB	₹2,000	A · 9/10	Rock-solid Very stable network, Mumbai DC. Backups cost extra
Hostinger VPS global DC	Unmanaged VPS	KVM 2 · 2vCPU / 8GB	₹550–700	A– · 8/10	Cheap long-term 8GB for the price, Docker templates. Needs 12–24mo commit for best rate.
MilesWeb / YouStable Mumbai DC · INR	Indian VPS	NVMe · 2vCPU / 4GB	₹600–900	A– · 8/10	UPI + local INR billing, Indian-hours support. Renewal > promo price.
AIC Cloud India DC · INR	Indian VPS	4GB plan (not the ₹99)	₹500–800	B · 7/10	Budget Cheapest INR entry & UPI, but smaller infra — verify uptime/backups.
Railway PaaS · US/EU	Container PaaS	Service + Volume · ~4GB	₹900–1,800+	B– · 6/10	Attach a Volume! Zero-ops & Mongo on one dashboard, but you MUST mount a persistent volume for the session, and billing is metered.

Render PaaS · US/EU	Container PaaS	Web Svc + Persistent Disk	₹600 + disk	C+ · 5/10	Disk required Easy GitHub deploys, but needs a Dockerfile for Chromium AND a paid persistent disk for the session.
DO App Platform PaaS	Managed App	Basic container	₹450–1,000	C · 4/10	Wrong shape Ephemeral FS + limited low-level Chrome control = fights a WhatsApp Web bot. Use a Droplet instead.

* INR ≈ at ₹84/\$; VPS shown at the 4GB tier your stack needs (not the cheapest 1GB teaser). † Railway metered – depends on actual RAM-hours + egress.

STEP 4 – THE ONE TRAP TO AVOID

Why "just git-push to Render/Railway" bites this app

! Ephemeral filesystem = QR re-scan on every deploy

On a normal PaaS, each deploy starts a fresh container with a blank disk. Your `.webjs_auth/` login is wiped, so the bot boots logged-out and waits for you to scan the WhatsApp QR again — mid-business-day, losing incoming messages. The fixes (Render persistent disk, Railway volume) work but add cost, force single-instance anyway, and still need a custom Dockerfile to install Chromium's system libraries. On a plain VPS none of this exists — the disk just persists. That's why the VPS options out-rank the "easier" PaaS ones *for this specific app*.

i MongoDB: run it local on the VPS, free

At 1,000 convos/day your sessions + leads + customers data is small. On a 4GB VPS, run MongoDB locally (free, ~300MB RAM) — no separate bill. Only reach for MongoDB Atlas (paid, from ~₹800/mo) if you go the PaaS route where local Mongo isn't practical. Your app already falls back to JSON files if Mongo is absent, so this is low-risk either way.

RECOMMENDATION

DO THIS

Hetzner CX22 (₹430/mo) — or DigitalOcean Bangalore if you want India-local

For pure value, nothing beats a **Hetzner CX22**: 2 vCPU / 4GB is exactly Chrome-plus-Mongo sized, the disk persists your WhatsApp session for free, and EU latency is irrelevant because your LLM lives in the US. If you specifically want a **Bangalore/Mumbai datacenter, UPI billing, and Indian-hours support**, step up to a **DigitalOcean 4GB Droplet (~₹2,000)** or a **MilesWeb/AIC 4GB Indian VPS (~₹600–900)** — same architecture, just local. Avoid the 1GB budget plans and the ephemeral-FS PaaS tier for this particular bot.

- ▶ Provision Ubuntu 22.04, 4GB VPS · install Node 20 + Chromium system libs
- ▶ Install MongoDB locally · set MONGODB_URI in .env
- ▶ Run under PM2 (pm2 start · pm2 save · pm2 startup) so it survives reboots
- ▶ Open /admin, scan the WhatsApp QR once — .webjs_auth/ persists it forever
- ▶ Set up a nightly snapshot/backup (₹80–100/mo) — cheap insurance

Report generated for Theaurax.in · stack-specific hosting analysis · costs indicative, verify live pricing before purchase.